



Fully Convolutional Neural Networks

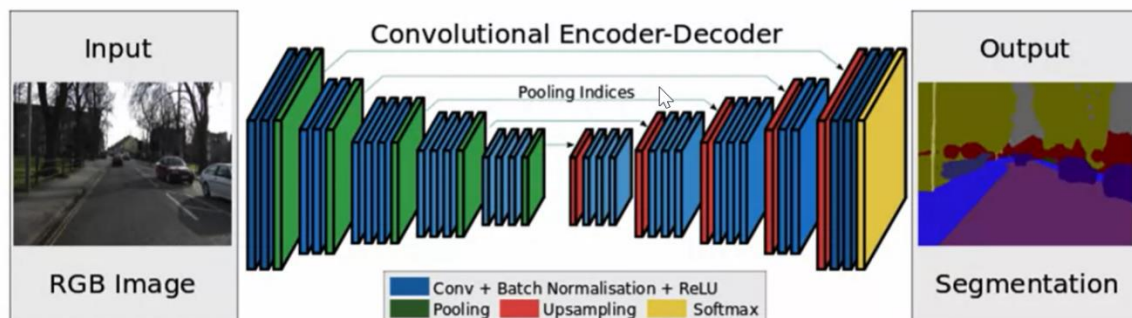


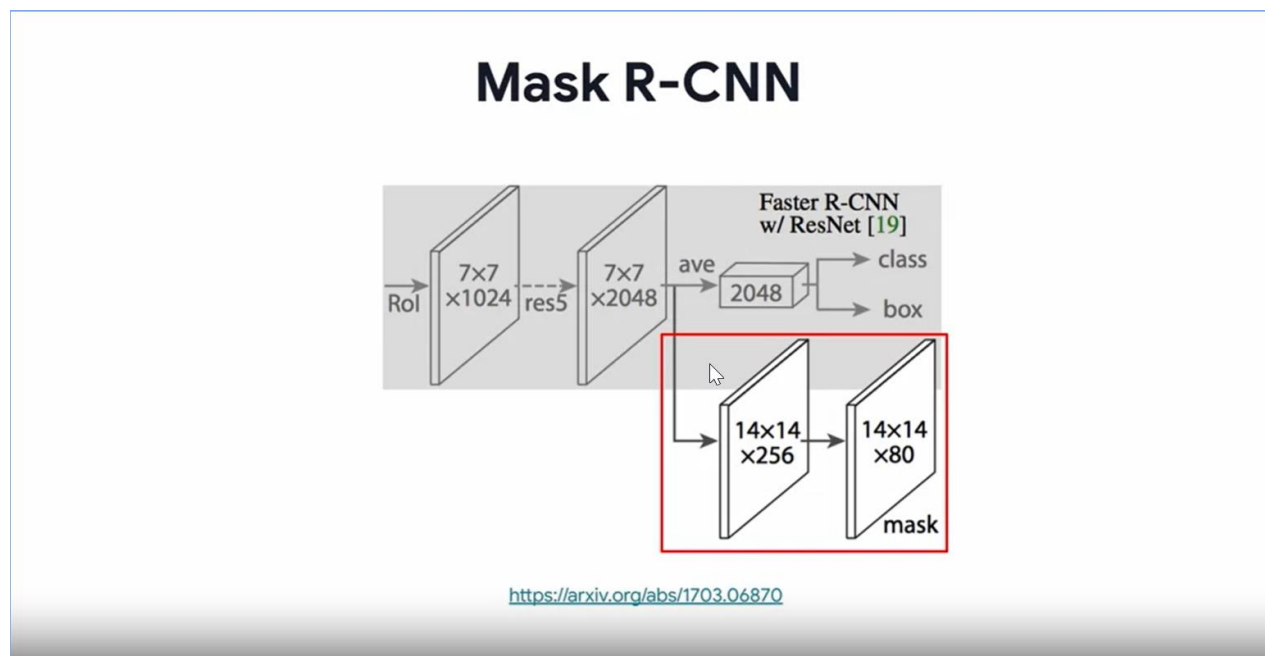

https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf

Comparison of Different FCNs




SegNet





FCN Architecture Details

Fully Convolutional Neural Networks





Fully Convolutional Networks for Semantic Segmentation
By Jonathan Long, Evan Shelhamer, Trevor Darrell
<https://arxiv.org/pdf/1411.4038.pdf>

Encoders


- Popular encoder architectures:
 - VGG-16
 - ResNet-50
 - MobileNet
- Reuse convolutional layers for feature extraction.
 - Do not reuse fully connected layers

Decoders

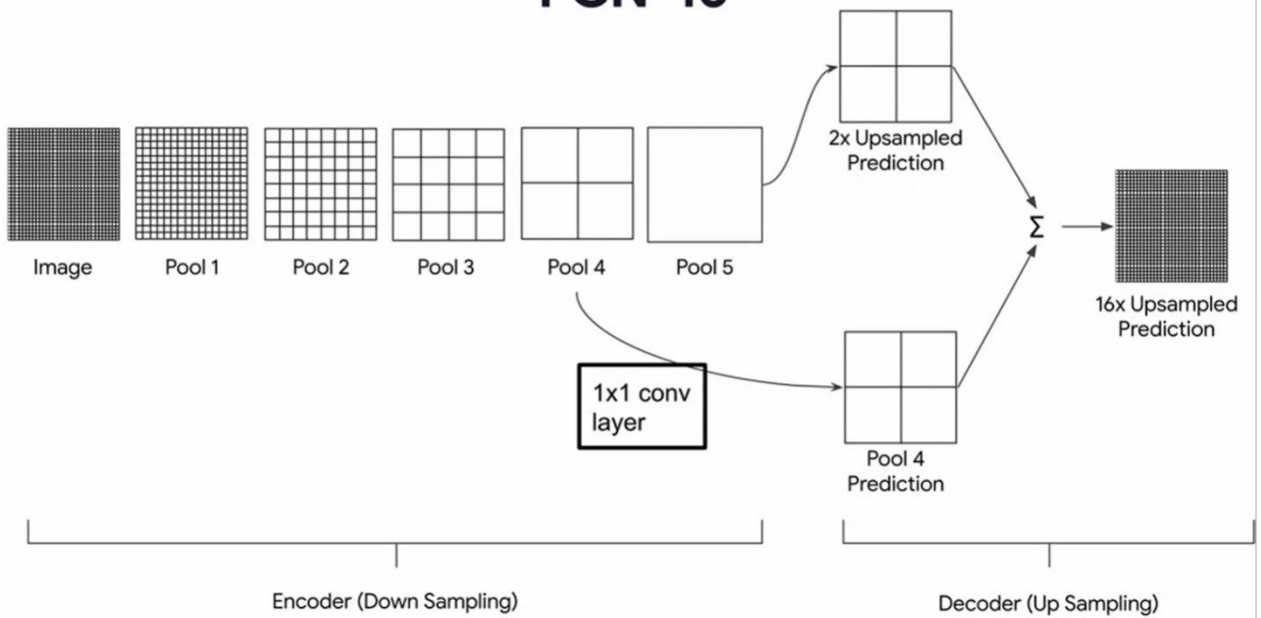



Height and width halved

FCN-32



FCN-16





Upsampling

- Upsampling is increasing height and width of the feature map.
- Two types of layers used in TensorFlow:
 - Simple Scaling - UpSampling2D
 - Transposed Convolution(Deconvolution) - Conv2DTranspose

Simple Scaling - UpSampling2D

- Upsampling2D scales up the image
- Two Types of scaling:
 - Nearest
 - Copies value from nearest pixel.
 - Bilinear
 - linear interpolation from nearby pixels.

UpSampling2D - Usage


```
x = UpSampling2D(  
    size=(2, 2),  
    data_format=None,  
    interpolation='nearest')(x)
```

size: int or tuple of two ints

data_format: 'channels_first', 'channels_last' or None

interpolation: 'nearest' or 'bilinear'

Transposed Convolution



Convolution->



<- Transposed Convolution


Conv2DTranspose

- Reverse of Convolution.
- Applied to output of a convolution operation.
- Uses a kernel of a specified size and stride in order to recreate the original input before the convolution operation.

```
Conv2DTranspose(  
    filters=32,  
    kernel_size=(3, 3)  
)
```


Encoder in Code

- The Cambridge Driving, labelled video database (aka CamVid) contains 10 minutes of 30fps video, segmented and labelled with 32 classes
- GitHub account [divamgupta](#) has taken a subsample of the CamVid dataset to create a smaller dataset.



conv1_1
conv1_2
pooling2
conv2_1
conv2_2
pooling2
conv3_1
conv3_2
conv3_3
pooling3
conv4_1
conv4_2
conv4_3
pooling4
conv5_1
conv5_2
conv5_3
pooling5

- [Divam Gupta](#)'s GitHub account containing a subsample of the CamVid dataset to create a smaller dataset.

```
def block(x, n_convs, filters, kernel_size, activation, pool_size, pool_stride, block_name):  
    for i in range(n_convs):  
        x = tf.keras.layers.Conv2D(filters=filters,  
                                    kernel_size=kernel_size, activation=activation,  
                                    padding='same',  
                                    name="{}_conv{}".format(block_name, i + 1))(x)  
  
    x = tf.keras.layers.MaxPooling2D(pool_size=pool_size, strides=pool_stride,  
                                     name="{}_pool{}".format(block_name, i+1 ))(x)  
  
    return x
```

conv4_1
conv4_2
conv4_3
pooling4

```
def VGG_16(image_input):
```

```
    x = block(image_input, n_convs=2, filters=64, kernel_size=(3,3),
               activation='relu', pool_size=(2,2), pool_stride=(2,2),
               block_name='block1')
```

```
    p1= x
```

```
    x = block(x, n_convs=2, filters=128, kernel_size=(3,3),
               activation='relu', pool_size=(2,2), pool_stride=(2,2),
               block_name='block2')
```

```
    p2 = x
```

```
    ...
```

conv1_1
conv1_2
pooling2

conv2_1
conv2_2
pooling2

```
    x = block(x, n_convs=3, filters=256, kernel_size=(3,3), activation='relu', pool_size=(2,2),
               pool_stride=(2,2), block_name='block3')
```

```
    p3 = x
```

conv3_1
conv3_2
conv3_3
pooling3

```
    x = block(x, n_convs=3, filters=512, kernel_size=(3,3), activation='relu', pool_size=(2,2),
               pool_stride=(2,2), block_name='block4')
```

```
    p4 = x
```

conv4_1
conv4_2
conv4_3
pooling4

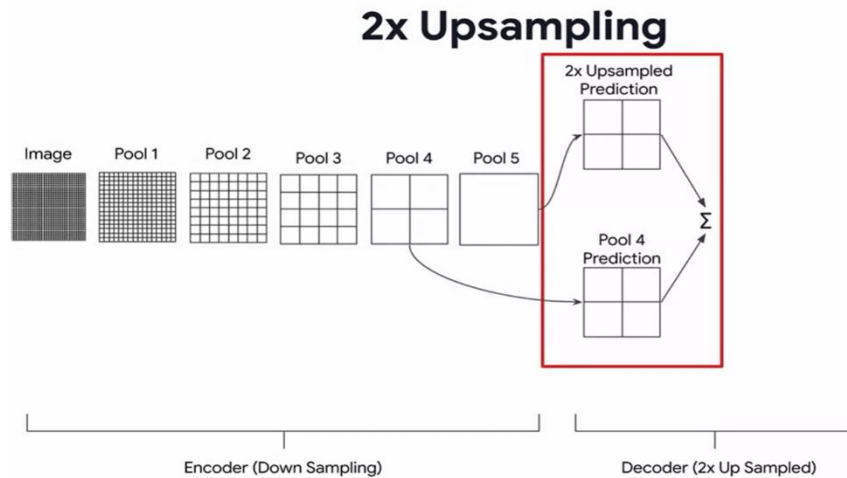
```
    x = block(x, n_convs=3, filters=512, kernel_size=(3,3), activation='relu', pool_size=(2,2),
               pool_stride=(2,2), block_name='block5')
```

```
    p5 = x
```

conv5_1
conv5_2
conv5_3
pooling5

```
    ..
```

Decoder in Code (FCN8)



Define Decoder - 2x UpSampling

```
def fcn8_decoder(convs, n_classes):
    f1, f2, f3, f4, f5 = convs

    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(4,4),
                                         strides=(2,2), use_bias=False)(f5)

    o = tf.keras.layers.Cropping2D(cropping=(1,1))(o)

    o2 = f4
    o2 = (tf.keras.layers.Conv2D(n_classes, (1,1),
                                  activation='relu', padding='same'))(o2)


    o = tf.keras.layers.Add()([o, o2])
    ...
```

1 x 1 Convolutions

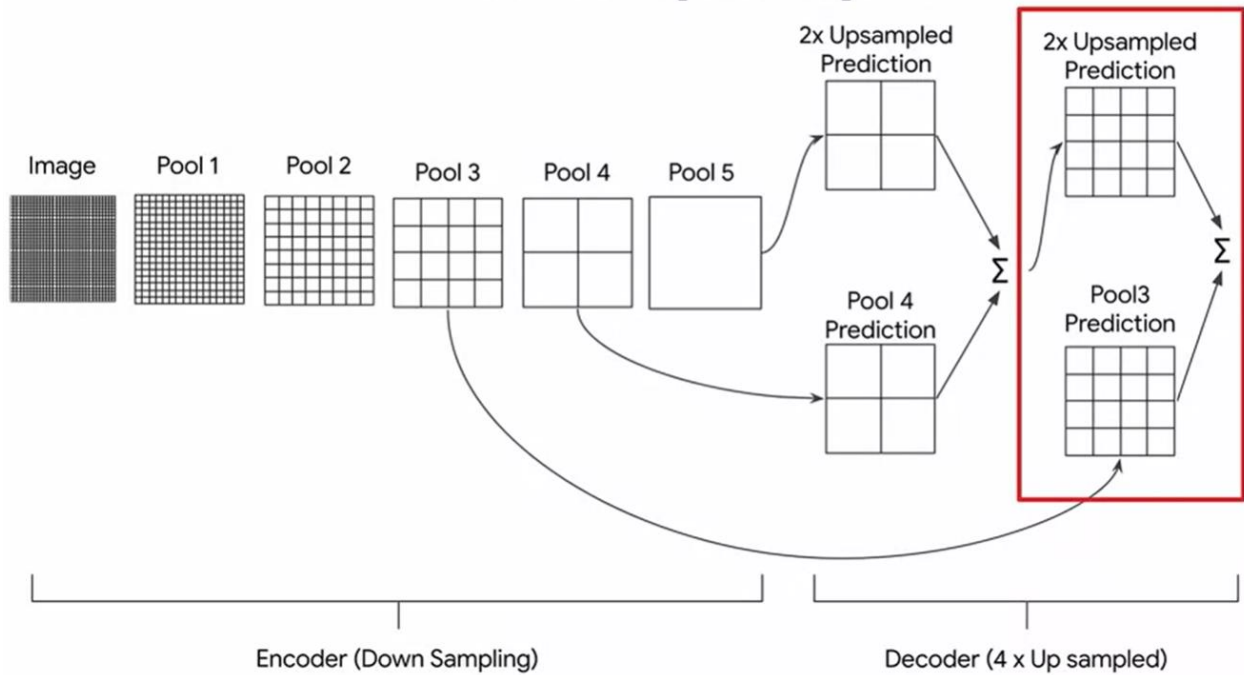
(B, F, H, W) - B = # batches; F = # filters, H, W = Height/Width

Apply a layer with N 1x1 Convolutions with stride of 1:

(B, N, H, W) - B = # batches; N = # filters, H, W = Height/Width



2 x 2 Upsampled



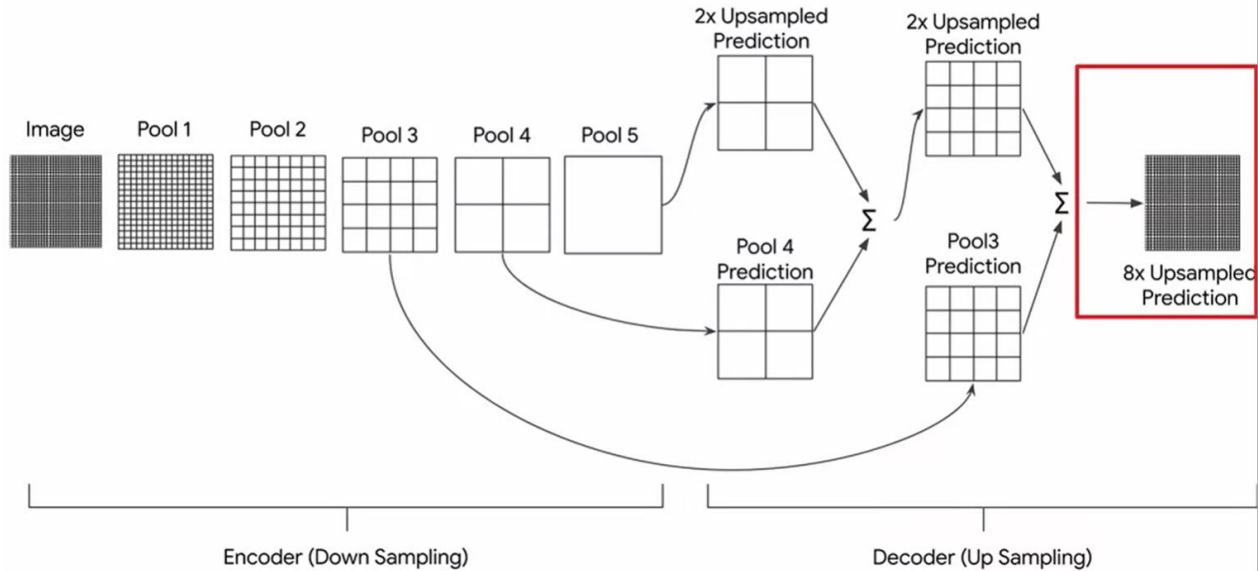
```
def fcn8_decoder(convs, n_classes):
    ...
    o = (tf.keras.layers.Conv2DTranspose( n_classes, kernel_size=(4,4),
                                           strides=(2,2)))(o)

    o = tf.keras.layers.Cropping2D(cropping=(1, 1))(o)

    o2 = ( tf.keras.layers.Conv2D(n_classes,(1,1), activation='relu',
                                   padding='same'))(f3)

    o = tf.keras.layers.Add()( [o, o2])
```

2 x 2 x 8 Up Sampled



Define Decoder

```
def fcn8_decoder(convs, n_classes):
    ...

    o = tf.keras.layers.Conv2DTranspose(n_classes, kernel_size=(8,8),
                                         strides=(8,8))(o)

    o = (tf.keras.layers.Activation('softmax'))(o)

    return o
```


Define Final Model

```
def segmentation_model():
    inputs = tf.keras.layers.Input(shape=(224,224,3,))
    convs = VGG_16(image_input=inputs)
    outputs = fcn8_decoder(convs, 12)
    model = tf.keras.Model(inputs=inputs, outputs=outputs)
    return model


model = segmentation_model()
```

Evaluation with IoU and Dice Score

Sample Visualization of Predicted Segments



Original Image




Predicted Segments




Ground Truth Segments

Area of Overlap


Area of Overlap = $\text{sum}(\text{True Positives})$



Combined Area = Total Pixels in predicted segmentation mask + Total Pixels in True Segmentation mask



Area of Union = Total Pixels in predicted segmentation mask + Total Pixels in True Segmentation mask - Area of Overlap




Calculate Areas

```
def class_wise_metrics(y_true, y_pred):
    ...
    smoothening_factor = 0.00001

    for i in range(n_classes):
        intersection = np.sum((y_pred == i) * (y_true == i))
        y_true_area = np.sum((y_true == i))
        y_pred_area = np.sum((y_pred == i))
        combined_area = y_true_area + y_pred_area
        union_area = combined_area - intersection
```

Intersection Over Union

$$\text{IOU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$



Calculate IOU

```
def class_wise_metrics(y_true, y_pred):
    class_wise_iou = []


    ...

    for i in range(n_classes):
        ...
        iou = (intersection) / (union_area)
        class_wise_iou.append(iou)

    return class_wise_iou
```

IOU Results

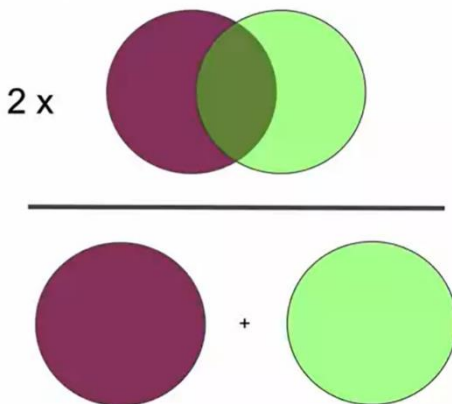
sky	0.8779669959482955
building	0.7570989578412737
column/pole	4.57875457665808e-10
road	0.915543155822588
side walk	0.7235628237658467
vegetation	0.7664541807647628
traffic light	3.0202657798187055e-05
fence	0.006380242448568188
vehicle	0.2950299461448835
pedestrian	0.0001264333276608086
bicyclist	0.023621930993270864
void	0.16456276759816527



You'll also use the dice score as one of the evaluation metrics in the Colab. The dice score is twice the area of overlap divided by the combined area. It can be used in similar circumstances to the intersection over union score, and they're often both used. The subtle difference between them is that the dice score tends to veer towards the average performance. Whereas the IOU helps you understand worst case performance. Code for this calculated class wise is also pretty straightforward. Iterate through the classes, calculate the intersection and the combined area for that class, and then calculate the score according to the formula. Again, we can get a metric for our classification using it. As you can see here, we did very well in some classes and we did quite poorly in others. Particularly alarming is that the pedestrian and cyclist scores are very poor.

You would not want to use a model like this in production. Hopefully, that was a useful introduction into how the fully convolutional layers architecture can be used for image segmentation. Now there's a lot to cover and there's no better way to do that than to dig into the code. Next, you're going to practice everything that you've seen so far. Building an image segmentation model to segment dashcam views for self-driving car tasks. You're going to use VGG16 for the encoder and FCN8 for the decoder, and you'll also use things like IOU and the dice score to evaluate your model's performance. In the next lesson, we'll switch gears and we'll look at the unit architecture for image segmentation

Dice Score

$$\text{Dice Score} = 2 \times \frac{\text{Area of Overlap}}{\text{Combined Area}}$$


Calculate Dice Score

```
def class_wise_metrics(y_true, y_pred):
    class_wise_dice_score = []

    ...

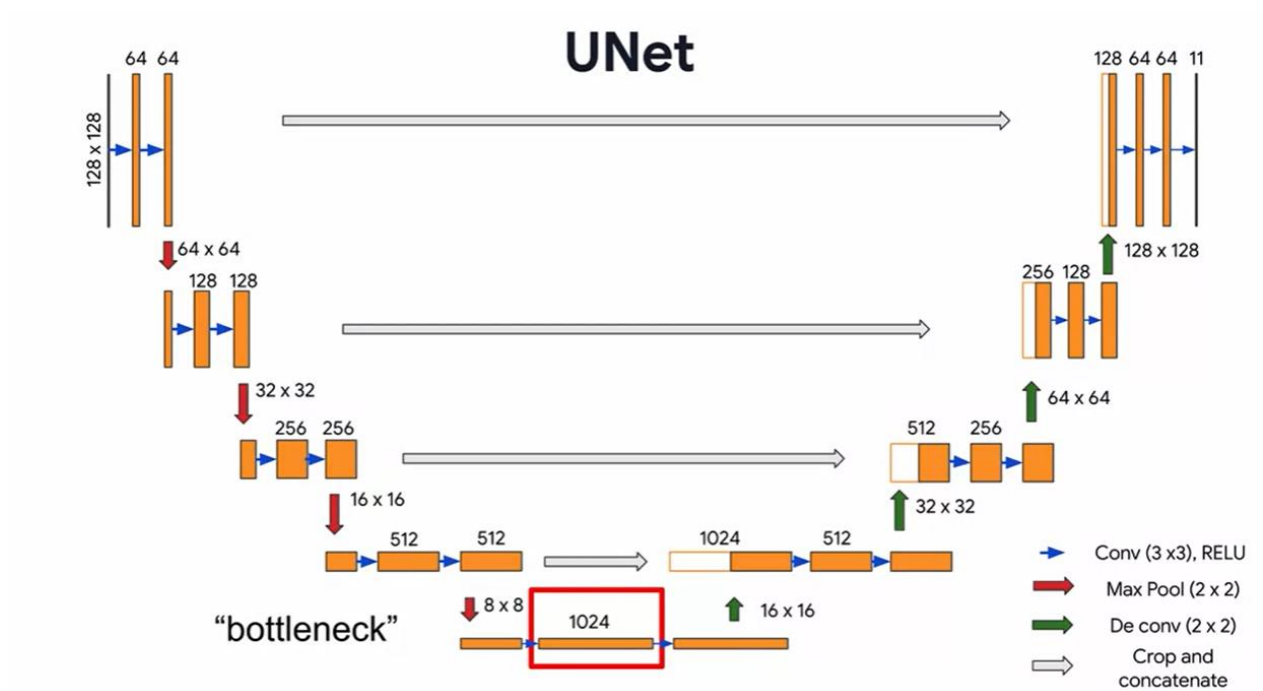
    for i in range(n_classes):
        ...
        dice_score = 2 * (intersection) / (combined_area)
        class_wise_dice_score.append(dice_score)

    return class_wise_dice_score
```


Dice Score Results

sky	0.9350185576821789
building	0.861760180856767
column/pole	9.15750915331616e-10
road	0.9559097147402678
side walk	0.8396129387346395
vegetation	0.8677883515092748
traffic light	6.040349126864078e-05
fence	0.012679586065167121
vehicle	0.45563416820437186
pedestrian	0.0002528346886971326
bicyclist	0.04615362426586659
void	0.2826172572009191





An additional element in the unit architecture, is a simple convolutional layer, which can further extract features, but it doesn't have a pooling layer to follow it. This convolutional layer is named the bottleneck. Data flows through that, before we get to the decoder side of the net.



And that's here and this is the right hand side of the U shape. Starting at the bottom, the 5th level with the 8 by 8 blocks gets up samples to 16 by 16. So, now were moving up to the 4th level of the U shape, but notice what happens next.

You'll take the 512 filters from the layer of the encoder that's at the same level as this decoder layer, in this case it's the 4th level. Notice that since the encoder layer and decoder layer are the same level in the unit, they also have the same height and width of 16 by 16 and they also have the same number of filters at 512 each.

So, you'll then concatenate the filters from the encoder, with the filters of the decoder, for a total of 1024 filters. You'll then pass this concatenated set of 1024 filters, through 2 convolutional layers.

And this pattern continues through the decoder, you'll up sample the blocks to 32 by 32 and move up to level 3. You'll take the filters from the encoder on the same level, concatenate them to the blocks from the decoder and pass the entire thing through 2 convolutional layers. And of course this works as well as up sampling the image, it will get fed through two layers of convolutional filters, that match the one on the previous level. But with the additional ones that are the skip over, from the equivalent filters on the encoder side.

So up sampled to 64 by 64, move up to level 2, combine the filters from the encoder with a decoder and pass them through two convolution layers. And finally, up sample to 128 by 128 and move up to level 1, concatenate the filters from the encoder and decoder and pass them through the two convolution layers.

